

```
#installing kaggle lib
! pip install kaggle
```

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.10/dist-packages (1.6.14)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.10/dist-packages (from kaggle) (1.16.0)
Requirement already satisfied: certifi>=2023.7.22 in /usr/local/lib/python3.10/dist-packages (from kaggle) (2024.6.2)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.8.2)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.31.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from kaggle) (4.66.4)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.10/dist-packages (from kaggle) (8.0.4)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-packages (from kaggle) (2.0.7)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from kaggle) (6.1.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->kaggle) (0.5.1)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.10/dist-packages (from python-slugify->kaggle) (1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->kaggle) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->kaggle) (3.7)
```

upload your kaggle.json file

```
#configuring the kaggle file path
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json
```

importing twitter sentiment dataset

```
# api to fetch the dataset from kaggle
!kaggle datasets download -d kazanova/sentiment140
```

```
Dataset URL: https://www.kaggle.com/datasets/kazanova/sentiment140
License(s): other
Downloading sentiment140.zip to /content
 96% 78.0M/80.9M [00:00<00:00, 125MB/s]
100% 80.9M/80.9M [00:00<00:00, 114MB/s]
```

```
# extracting the compressed dataset
from zipfile import ZipFile
dataset = '/content/sentiment140.zip'
```

```
with ZipFile(dataset, 'r') as zip:
    zip.extractall()
    print('The dataset is extracted')
```

```
The dataset is extracted
```

importing the dependencies

```
import numpy as np
import pandas as pd
import re
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
import nltk
nltk.download('stopwords')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
True
```

```
#printing the stopwords in english
print(stopwords.words('english'))
```

```
['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'your', 'yours', 'yourself',
```

data processing

```
# loading the data from csv file to pandas dataframe
twitter_data = pd.read_csv('/content/training.1600000.processed.noemoticon.csv', encoding = 'ISO-8859-1')
```

```
# checking the num of rows and cols
twitter_data.shape
```

```
(1599999, 6)
```

```
# printing the first 5 rows of the dataframe
twitter_data.head()
```

```

0    1467810369    Mon Apr 06 22:19:45 PDT 2009    NO_QUERY    _TheSpecialOne_    @switchfoot
http://twitpic.com/2y1zl - Awww, that's a bummer. You shoulda got David Carr of Third Day to do it. ;D

0    0    1467810672    Mon Apr 06 22:19:49 PDT 2009    NO_QUERY    scotthamilton    is upset that he can't update his Facebook by ...

1    0    1467810917    Mon Apr 06 22:19:53    NO_QUERY    mattycus    @Kenichan I dived many times for the ball Man

```

```
# naming the cols and reading the dataset again
```

```
column_names = ['target', 'id', 'date', 'flag', 'user', 'text']
twitter_data = pd.read_csv('/content/training.1600000.processed.noemoticon.csv', names=column_names, encoding = 'ISO-8859-1')
```

```
# checking the num of rows and cols
twitter_data.shape
```

```
(1600000, 6)
```

```
# printing the first 5 rows of the dataframe
twitter_data.head()
```

```

target    id    date    flag    user    text
0    0    1467810369    Mon Apr 06 22:19:45 PDT 2009    NO_QUERY    _TheSpecialOne_    @switchfoot
http://twitpic.com/2y1zl - Awww, t...

1    0    1467810672    Mon Apr 06 22:19:49 PDT 2009    NO_QUERY    scotthamilton    is upset that he can't update his Facebook by ...

```

```
# counting the number of missing values in the dataset
twitter_data.isnull().sum()
```

```

target    0
id        0
date      0
flag      0
user      0
text      0
dtype: int64

```

```
# checking the distribution of the target col
twitter_data['target'].value_counts()
```

```

target
0    800000
4    800000
Name: count, dtype: int64

```

converting the target value '4' as '1'

```
twitter_data.replace({'target':{4:1}}, inplace=True)
```

0--> negative tweet

1--> postive tweet

checking the distribution of the target col

```
twitter_data['target'].value_counts()
```

```
target
0    800000
1    800000
Name: count, dtype: int64
```

stemming

stemming is the process of reducing a word to its root word

eg : actor, actress, acting = act

```
port_stem = PorterStemmer()
```

```
def stemming(content):
```

```
    stemmed_content = re.sub('[^a-zA-Z]', ' ', content)
    stemmed_content = stemmed_content.lower()
    stemmed_content = stemmed_content.split()
    stemmed_content = [port_stem.stem(word) for word in stemmed_content if not word in stopwords.words('english')]
    stemmed_content = ' '.join(stemmed_content)
```

```
    return stemmed_content
```

```
twitter_data['stemmed_content'] = twitter_data['text'].apply(stemming)
```

printing the first 5 rows of the dataframe

```
twitter_data.head()
```

```
target    id    date    flag    user    text    stemmec
0         0  1467810369  22:19:45 NO_QUERY _TheSpecialOne_ http://twitpic.com/2y1zl - Awww, t...
Mon Apr 06 PDT 2009
0         0  1467810369  22:19:45 NO_QUERY _TheSpecialOne_ http://twitpic.com/2y1zl - Awww, t...
Mon Apr 06 PDT 2009
0         0  1467810369  22:19:45 NO_QUERY _TheSpecialOne_ http://twitpic.com/2y1zl - Awww, t...
Mon Apr 06 PDT 2009
0         0  1467810369  22:19:45 NO_QUERY _TheSpecialOne_ http://twitpic.com/2y1zl - Awww, t...
Mon Apr 06 PDT 2009
```

```
print(twitter_data['stemmed_content'])
```

```
0    switchfoot http twitpic com z1 awww bummer sho...
1    upset updat facebook text might cri result sch...
2    kenichan dive mani time ball manag save rest g...
3    whole bodi feel itchi like fire
4    nationwideclass behav mad see
...
1599995    woke school best feel ever
1599996    thewdb com cool hear old walt interview http b...
1599997    readi mojo makeov ask detail
1599998    happi th birthday boo alll time tupac amaru sh...
1599999    happi charitytuesday thenspcc sparkschar speak...
Name: stemmed_content, Length: 1600000, dtype: object
```

```
print(twitter_data['target'])
```

```
0    0
1    0
2    0
3    0
```

```

4      0
      ..
1599995 1
1599996 1
1599997 1
1599998 1
1599999 1
Name: target, Length: 1600000, dtype: int64

```

```

# separating the data and label
x = twitter_data['stemmed_content'].values
y = twitter_data['target'].values

```

+ Code + Text

```
print(x)
```

```

['switchfoot http twitpic com zl awww bummer shoulda got david carr third day'
 'upset updat facebook text might cri result school today also blah'
 'kenichan dive mani time ball manag save rest go bound' ...
 'readi mojo makeov ask detail'
 'happi th birthday boo alll time tupac amaru shakur'
 'happi charitytuesday thenspcc sparkshar speakingup h']

```

```
print(y)
```

```
[0 0 0 ... 1 1 1]
```

splitting the training data and the testing data

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, stratify=y, random_state=2)
```

```
print(x.shape, x_train.shape, x_test.shape)
```

```
(1600000,) (1280000,) (320000,)
```

```
print(x_train)
```

```

['watch saw iv drink lil wine' 'hatermagazin'
 'even though favourit drink think vodka coke wipe mind time think im gonna find new drink'
 ... 'eager monday afternoon'
 'hope everyon mother great day wait hear guy store tomorrow'
 'love wake folger bad voic deeper']

```

```
print(x_test)
```

```

['mmangen fine much time chat twitter hubbi back summer amp tend domin free time'
 'ah may show w ruth kim amp geoffrey sanhueza'
 'ishatara mayb bay area thang dammit' ...
 'destini nevertheless hooray member wonder safe trip' 'feel well'
 'supersandro thank']

```

converting the textual data to numerical data

```
vectorizer = TfidfVectorizer()
```

```
x_train = vectorizer.fit_transform(x_train)
```

```
x_test = vectorizer.transform(x_test)
```

```
print(x_train)
```

```

(0, 443066) 0.4484755317023172
(0, 235045) 0.41996827700291095
(0, 109306) 0.3753708587402299
(0, 185193) 0.5277679060576009
(0, 354543) 0.3588091611460021
(0, 436713) 0.27259876264838384
(1, 160636) 1.0
(2, 288470) 0.16786949597862733
(2, 132311) 0.2028971570399794
(2, 150715) 0.18803850583207948
(2, 178061) 0.1619010109445149
(2, 409143) 0.15169282335109835
(2, 266729) 0.24123230668976975
(2, 443430) 0.3348599670252845
(2, 77929) 0.31284080750346344

```

```

(2, 433560)    0.3296595898028565
(2, 406399)    0.32105459490875526
(2, 129411)    0.29074192727957143
(2, 407301)    0.18709338684973031
(2, 124484)    0.1892155960801415
(2, 109306)    0.4591176413728317
(3, 172421)    0.37464146922154384
(3, 411528)    0.27089772444087873
(3, 388626)    0.3940776331458846
(3, 56476)     0.5200465453608686
:
(1279996, 390130) 0.22064742191076112
(1279996, 434014) 0.2718945052332447
(1279996, 318303) 0.21254698865277746
(1279996, 237899) 0.2236567560099234
(1279996, 291078) 0.17981734369155505
(1279996, 412553) 0.18967045002348676
(1279997, 112591) 0.7574829183045267
(1279997, 273084) 0.4353549002982409
(1279997, 5685)   0.48650358607431304
(1279998, 385313) 0.4103285865588191
(1279998, 275288) 0.38703346602729577
(1279998, 162047) 0.34691726958159064
(1279998, 156297) 0.3137096161546449
(1279998, 153281) 0.28378968751027456
(1279998, 435463) 0.2851807874350361
(1279998, 124765) 0.32241752985927996
(1279998, 169461) 0.2659980990397061
(1279998, 93795)  0.21717768937055476
(1279998, 412553) 0.2816582375021589
(1279999, 96224)  0.5416162421321443
(1279999, 135384) 0.6130934129868719
(1279999, 433612) 0.3607341026233411
(1279999, 435572) 0.31691096877786484
(1279999, 31410)  0.248792678366695
(1279999, 242268) 0.19572649660865402

```

```
print(x_test)
```



```

(0, 420984)    0.17915624523539803
(0, 409143)    0.31430470598079707
(0, 398906)    0.3491043873264267
(0, 388348)    0.21985076072061738
(0, 279082)    0.1782518010910344
(0, 271016)    0.4535662391658828
(0, 171378)    0.2805816206356073
(0, 138164)    0.23688292264071403
(0, 132364)    0.25525488955578596
(0, 106069)    0.3655545001090455
(0, 67828)     0.26800375270827315
(0, 31168)     0.16247724180521766
(0, 15110)     0.1719352837797837
(1, 366203)    0.24595562404108307
(1, 348135)    0.4739279595416274
(1, 256777)    0.28751585696559306
(1, 217562)    0.40288153995289894
(1, 145393)    0.575262969264869
(1, 15110)     0.211037449588008
(1, 6463)      0.30733520460524466
(2, 400621)    0.4317732461913093
(2, 256834)    0.2564939661498776
(2, 183312)    0.5892069252021465
(2, 89448)     0.36340369428387626
(2, 34401)     0.37916255084357414
:
(319994, 123278) 0.4530341382559843
(319995, 444934) 0.3211092817599261
(319995, 420984) 0.22631428606830145
(319995, 416257) 0.23816465111736276
(319995, 324496) 0.3613167933647574
(319995, 315813) 0.28482299145634127
(319995, 296662) 0.39924856793840147
(319995, 232891) 0.25741278545890767
(319995, 213324) 0.2683969144317078
(319995, 155493) 0.2770682832971668
(319995, 109379) 0.30208964848908326
(319995, 107868) 0.3339934973754696
(319996, 438709) 0.4143006291901984
(319996, 397506) 0.9101400928717545
(319997, 444770) 0.2668297951055569
(319997, 416695) 0.29458327588067873
(319997, 349904) 0.32484594100566083
(319997, 288421) 0.48498483387153407

```

```
(319997, 261286)      0.37323893626855326
(319997, 169411)      0.403381646999604
(319997, 98792)       0.4463892055808332
(319998, 438748)      0.719789181620468
(319998, 130192)      0.6941927210956169
(319999, 400636)      0.2874420848216212
(319999, 389755)      0.9577980203954275
```

training ml model

Logistic Regression

```
model = LogisticRegression(max_iter = 1000)
```

```
model.fit(x_train, y_train)
```

```
LogisticRegression
LogisticRegression(max_iter=1000)
```

model evaluation

accuracy score

```
# accuracy score on the training data
x_train_prediction = model.predict(x_train)
training_data_accuracy = accuracy_score(y_train, x_train_prediction)
```

```
print('Accuracy score on the training data : ', training_data_accuracy)
```

```
Accuracy score on the training data : 0.81018984375
```

```
# accuracy score on the test data
x_test_prediction = model.predict(x_test)
test_data_accuracy = accuracy_score(y_test, x_test_prediction)
```

```
print('Accuracy score on the test data : ', test_data_accuracy)
```

```
Accuracy score on the test data : 0.7780375
```

model accuracy score = 77.8%

saving the trained model

```
import pickle
```

```
filename = 'trained_model.sav'
pickle.dump(model, open(filename, 'wb'))
```


using the saved model for future predictions

```
# loading the saved model
loaded_model = pickle.load(open('/content/trained_model.sav', 'rb'))
```

```
x_new = ["sad"]
x_new = vectorizer.transform(x_new)
```

```
prediction = loaded_model.predict(x_new)
#print(prediction)
```

```
if(prediction[0] == 0):
    print('negative tweet')
else:
    print('positive tweet')
```

 [0]
negative tweet